

Revisiting the Student Sectioning Problem through Graph Representation Learning

Abstract

Student sectioning is a fundamental subproblem of university course timetabling that aims to partition students into balanced sections while preserving scheduling flexibility and efficient resource utilization. Classical clustering-based approaches typically represent students using binary student-course enrollment matrices, which may not adequately capture the relational structure induced by shared course participation.

This paper revisits student sectioning from the perspective of graph representation learning. We construct a weighted student co-enrollment graph after removing constant features, and employ Node2Vec to learn low-dimensional student embeddings from biased random-walk contexts. The learned representations are then provided to standard clustering algorithms for section formation. To distinguish graph-based representation learning from dimensionality reduction and direct graph clustering, we compare the proposed approach with the conventional enrollment representation, PCA followed by KMeans, and spectral clustering on the co-enrollment graph.

Experiments on six real-world university courses involving 341 students show that Node2Vec achieves an average Silhouette Score of 0.611 with the neutral configuration ($p=1$, $q=1$), compared with 0.153 for the conventional representation, 0.460 for PCA followed by KMeans, and 0.213 for spectral clustering. The best Node2Vec configuration ($p=1$, $q=0.5$) achieves 0.622. Node2Vec outperforms PCA and the conventional representation on all six courses, with large effect sizes (Cliff's $\delta = 1.0$). However, after Holm correction for multiple comparisons, the improvements are not statistically significant at $\alpha = 0.05$ due to the small number of independent datasets ($n = 6$).

The results indicate that graph representation learning can capture useful structural information beyond conventional dimensionality reduction. The proposed approach can be incorporated into existing student sectioning pipelines without modifying the downstream clustering algorithms, providing a simple framework for improving student groupings used in subsequent timetabling processes.

Keywords: Student sectioning, Graph representation learning, Node2Vec, Educational data mining, Course timetabling, Clustering

1 Introduction

Student sectioning is a classical subproblem of university course timetabling, aiming to divide enrolled students into balanced instructional sections while preserving the flexibility required for subsequent scheduling and resource allocation. The quality of sectioning can affect downstream timetabling by influencing the compatibility of student groups with available classrooms, instructors, and time slots. Consequently, student sectioning has long been studied as an important problem within educational scheduling and constraint satisfaction.

Research on student sectioning dates back more than two decades. Early studies investigated clustering-based approaches, including fuzzy clustering with feature selection and later nature-inspired optimization methods, demonstrating the potential of data-driven grouping strategies for student sectioning [1–3]. Other studies formulated student sectioning as an optimization or constraint programming problem within larger university timetabling systems [4–8]. Together, these studies established important computational foundations for student sectioning and highlighted its role in educational resource management.

A common representation in clustering-based approaches is a student-course enrollment matrix, in which each student is represented by a binary vector indicating the courses in which the student is enrolled. This representation provides a direct description of student-course relationships and can be effective for sectioning. However, students can also be viewed as nodes in a relational network, where students are connected according to shared course enrollments. Such a representation makes it possible to model relationships between students not only through individual shared courses, but also through their connectivity patterns in the resulting co-enrollment network.

An important question is therefore whether the relational structure of student enrollments can provide useful information beyond the original enrollment representation. This question is particularly relevant because conventional dimensionality-reduction techniques can already substantially improve clustering in high-dimensional enrollment spaces. Thus, evaluating graph-based representations requires distinguishing their contribution from the general benefit of reducing the dimensionality of the input space.

Graph representation learning provides a natural framework for addressing this question. In particular, random-walk-based methods such as Node2Vec learn continuous node representations from the contexts generated by biased random walks on a graph [9, 10]. Node2Vec extends DeepWalk by introducing return and in-out parameters that control the balance between local and global graph exploration. These representations can encode structural relationships among nodes and provide low-dimensional features that can subsequently be used by conventional machine-learning algorithms.

Graph representation learning appears to have received limited attention in the student sectioning literature, despite the natural relational structure of student co-enrollment data. This motivates the central question addressed in this paper:

Can graph representation learning provide student representations that improve sectioning quality beyond both conventional enrollment representations and standard dimensionality reduction?

To investigate this question, we model student enrollments as a weighted co-enrollment graph and learn student embeddings using Node2Vec, a random-walk-based graph representation learning method that extends DeepWalk through biased random walks controlled by return and in-out parameters. The learned embeddings are then supplied to standard clustering algorithms for section formation. Importantly, the clustering stage is kept unchanged, allowing the effect of the learned representation to be examined independently of the choice of clustering algorithm. To assess whether the observed improvement is attributable merely to dimensionality reduction, we compare the proposed representation with a PCA-based representation having the same dimensionality. We also include spectral clustering as a graph-based baseline.

Experimental results on six real-world university courses involving 341 students show that Node2Vec provides substantially better clustering quality than the conventional enrollment representation and spectral clustering. PCA also produces a considerable improvement, but Node2Vec achieves higher average Silhouette Scores than PCA followed by KMeans, indicating that graph-based representations capture useful structural information beyond linear dimensionality reduction. Node2Vec achieves an average Silhouette Score of 0.611 with the neutral configuration ($p=1$, $q=1$), compared with 0.460 for PCA+KMeans. The best configuration ($p=1$, $q=0.5$) achieves 0.622. Additional experiments investigate parameter sensitivity, clustering stability, statistical significance, computational cost, and the resulting low-dimensional representations.

The main contributions of this paper are summarized as follows.

1. We revisit the classical student sectioning problem from the perspective of graph representation learning by constructing a weighted student co-enrollment graph and learning student representations using Node2Vec, a random-walk-based method with controllable exploration bias.
2. We systematically compare graph-based representations with both the conventional student-course enrollment representation and a PCA-based dimensionality-reduction baseline, thereby distinguishing the effect of graph representation learning from the general benefit

of dimensionality reduction. Spectral clustering is additionally evaluated as a direct graph-based clustering baseline.

3. We demonstrate on six real-world university courses that Node2Vec achieves substantially better clustering performance than the conventional approach and spectral clustering, with an average Silhouette Score of 0.611, compared with 0.460 for PCA followed by KMeans, 0.153 for the conventional enrollment representation, and 0.213 for spectral clustering. The improvement is consistent across all six courses, although after Holm correction the statistical significance does not reach $\alpha = 0.05$ due to the small number of independent datasets.
4. We provide comprehensive empirical analyses of Node2Vec parameter sensitivity, clustering stability across 20 random seeds, statistical significance with effect sizes, runtime, and embedding structure, providing a detailed characterization of the proposed framework.
5. We show that graph-based student representations can be incorporated into existing clustering-based sectioning pipelines without modifying the downstream clustering algorithms, providing a simple framework for investigating relational representations in educational scheduling.

The remainder of this paper is organized as follows. Section 2 reviews previous work on student sectioning and graph representation learning. Section 3 presents the proposed framework. Section 4 reports the experimental results and comparative analyses. Section 5 discusses the findings, while Section 6 concludes the paper.

2 Related Work

This work revisits the classical student sectioning problem from the perspective of graph representation learning. The related literature is therefore reviewed from two complementary directions: classical approaches to student sectioning and course timetabling, and graph representation learning for relational data. We then position the present study at the intersection of these two directions.

2.1 Student Sectioning and Course Timetabling

Student sectioning is a fundamental component of university course timetabling and educational scheduling. Its objective is to partition students into instructional sections while satisfying scheduling constraints and making efficient use of institutional resources. Because section assignments affect the flexibility of subsequent timetable construction, student sectioning has long been studied as an optimization problem within operations research and constraint programming.

Early studies primarily formulated student sectioning as part of larger timetabling systems. Aubin and Ferland [4] incorporated section generation into a university timetabling framework, while Hertz and Robert [5] decomposed the scheduling process into assignment subproblems including student sectioning. Carter [6] considered grouping students according to course selections before timetable construction with the objective of reducing subsequent scheduling conflicts. Other approaches employed tabu search [7] and priority-based constraint satisfaction techniques [11]. These studies established student sectioning as an important component of integrated university timetabling systems.

Data-driven approaches subsequently introduced clustering techniques for student sectioning. Fuzzy clustering with feature selection was proposed in [1, 2], representing students through enrollment-based feature vectors and applying fuzzy c-means clustering together with criteria for evaluating section balance. A Fish School clustering algorithm was later developed specifically for student sectioning, demonstrating the applicability of nature-inspired optimization to this problem [3]. Although these approaches differ in their optimization mechanisms, they generally represent students through their course enrollment information and subsequently apply a clustering or partitioning procedure.

2.2 Graph Representation Learning

Graph representation learning provides an alternative way of representing objects when their relationships are an important source of information. Instead of treating observations as independent feature vectors, graph embedding methods exploit graph topology to learn

continuous low-dimensional representations that reflect structural relationships among nodes.

DeepWalk [9] introduced a widely adopted random-walk-based approach to graph representation learning by generating truncated random walks and applying the Skip-Gram model to the resulting node sequences. The learned embeddings capture patterns of node co-occurrence induced by the graph structure. Node2Vec [10] subsequently extended this framework through biased random walks that provide greater control over local and global graph exploration. Related studies have investigated weighted random walks and other graph sampling strategies [12]. Graph representation and graph-based learning techniques have also been employed in a variety of application-oriented problems, including network embedding, structured optimization, and graph neural network-based anomaly detection [13–15].

These developments illustrate the broader applicability of graph-based learning beyond conventional graph benchmarks. In particular, they suggest that exploiting relational structure can be useful when the relationships among entities contain information that is difficult to capture directly from their original feature descriptions.

2.3 Graph Representation Learning for Student Sectioning

The student–course enrollment data used in sectioning naturally induces a relational structure among students. Two students can be connected according to the courses they jointly attend, resulting in a student co-enrollment graph. Such a graph provides an alternative representation of the same underlying enrollment information and makes it possible to investigate whether structural patterns in the co-enrollment network can improve student grouping.

To our knowledge, the use of graph representation learning as an explicit representation stage for student sectioning has received little attention. The present work investigates this intersection by constructing a weighted student co-enrollment graph and learning student embeddings using DeepWalk. The resulting embeddings are subsequently supplied to standard clustering algorithms, without modifying the clustering procedure itself.

An important aspect of the present study is that the improvement is not evaluated only against the original

high-dimensional enrollment representation. Since dimensionality reduction can itself substantially affect clustering performance, we additionally consider PCA followed by KMeans as a representation-based baseline with the same embedding dimensionality as DeepWalk. We also evaluate spectral clustering directly on the co-enrollment graph as a graph-based clustering baseline. These comparisons allow us to examine whether the observed benefit of DeepWalk is attributable merely to dimensionality reduction or whether graph-based representation learning provides additional useful information for student sectioning.

The present study therefore complements earlier clustering-based approaches to student sectioning rather than replacing them. Its central objective is to examine whether a Node2Vec-based representation learning stage can improve the quality of student representations while leaving the subsequent clustering pipeline unchanged.

3 Methodology

The proposed framework revisits the student sectioning problem from a graph representation learning perspective. Instead of representing students solely through a binary student–course enrollment matrix, we construct a student co-enrollment graph and learn continuous vector representations of students from the topology of this graph. These representations are subsequently used as input to conventional clustering algorithms to generate student sections.

The overall framework consists of four main stages: (1) construction of the student co-enrollment graph, (2) graph representation learning using random walks, (3) clustering of the learned student representations, and (4) evaluation of the resulting student sections. Figure 1 illustrates the overall workflow.

3.1 Student Co-enrollment Graph

Let $S = \{s_1, s_2, \dots, s_n\}$ denote the set of students and let $C = \{c_1, c_2, \dots, c_m\}$ denote the set of courses. Student enrollment information can be represented by a binary matrix $X \in \{0, 1\}^{n \times m}$, where

$$X_{ic} = \begin{cases} 1, & \text{if student } s_i \text{ is enrolled in course } c, \\ 0, & \text{otherwise.} \end{cases}$$




Fig. 1: Overview of the proposed framework. A student–course enrollment matrix is transformed into a student co-enrollment graph, from which random walks are generated and fed to a Skip-Gram (Word2Vec) model to learn low-dimensional student embeddings. The learned embeddings are then clustered to produce student sections.

Although the enrollment matrix provides a direct representation of student–course relationships, it does not explicitly represent relationships between students. We therefore transform the enrollment information into a student co-enrollment graph.

We define an undirected weighted graph $G = (V, E, W)$, where each vertex $v_i \in V$ corresponds to a student s_i . An edge $(v_i, v_j) \in E$ is established when students s_i and s_j share at least one enrolled course. The edge weight measures the strength of their co-enrollment relationship.

In the basic version of the proposed framework, the edge weight between students s_i and s_j is defined as the number of courses in which they are simultaneously enrolled:

$$w_{ij} = \sum_{c \in C} X_{ic} X_{jc}.$$

Consequently, students who share several courses are connected by stronger edges, whereas students who share only one course have a weaker connection. This weighted representation preserves more information than a binary adjacency matrix and provides a natural basis for subsequent random-walk-based representation learning.

The resulting graph can be represented by a weighted adjacency matrix $W \in \mathbb{R}^{n \times n}$, where $W_{ij} = w_{ij}$ for connected students and $W_{ij} = 0$ otherwise. Because the co-enrollment relationship is symmetric, the resulting graph is undirected and $W_{ij} = W_{ji}$.

3.2 Preprocessing: Removal of Constant Features

Before constructing any representation, we apply a standard preprocessing step to the student–course enrollment matrix. Courses that are taken by every student in a sectioning instance have zero variance and provide no discriminative information for clustering. Such constant features are therefore removed from the enrollment matrix before any representation or graph is constructed.

Specifically, for each target course, we identify courses c for which $X_{ic} = 1$ for all students s_i in the student set. These universally shared courses are removed, and the remaining enrollment matrix is used for all subsequent steps: BoW representation, PCA projection, graph construction, and random-walk-based embedding. This ensures that all methods operate on the same preprocessed data.

In the six datasets studied here, exactly one constant course is identified and removed per sectioning instance. The resulting graphs are non-complete, with average density approximately 0.72, in contrast to the fully connected graphs that arise when the constant course is included.

3.3 Random-walk-based Student Representation Learning

After constructing the co-enrollment graph, we learn a vector representation for each student using Node2Vec [10], a random-walk-based graph embedding method. The main idea is to treat the graph analogously to a corpus:

students correspond to words, and random walks correspond to sentences. Students that frequently occur in similar graph contexts are therefore expected to obtain similar vector representations.

Node2Vec generates biased random walks in which the transition probability depends on both the previous and the next-to-previous node in the walk. Let t denote the previous node and v denote the current node. The unnormalized transition probability to a neighbor x is defined as

$$\pi(v, x) = \alpha_{pq}(t, x) \cdot w_{vx},$$

where w_{vx} is the edge weight and the bias factor $\alpha_{pq}(t, x)$ is

$$\alpha_{pq}(t, x) = \begin{cases} 1/p, & \text{if } d(t, x) = 0, \\ 1, & \text{if } d(t, x) = 1, \\ 1/q, & \text{if } d(t, x) = 2, \end{cases}$$

where $d(t, x)$ denotes the shortest-path distance between nodes t and x . The parameter p controls the likelihood of returning to the previous node (return parameter), and q controls whether the walk explores locally (BFS-like, $q < 1$) or globally (DFS-like, $q > 1$). Setting $p = 1$ and $q = 1$ recovers the standard unbiased random walk used by DeepWalk.

Thus, students with stronger co-enrollment relationships have a higher probability of being visited during the random walk, and the bias parameters allow the walk to emphasize either local neighborhood structure or broader graph exploration. Repeating this procedure produces a collection of student sequences (Algorithm 1),

$$\mathcal{W} = \{W_1, W_2, \dots, W_R\},$$

where R denotes the total number of generated walks.

The resulting walks capture both local and higher-order relationships in the student graph (Figure 2 illustrates an example walk). For example, two students who do not directly share a course may nevertheless occur frequently in similar random-walk contexts because they are connected through common groups of students. Such relationships cannot be directly captured by treating each student’s enrollment vector independently.




Fig. 2: Example random walk on the student graph for Course 2, starting at node 40 and ending at node 45: $\{40, 36, 25, 12, 14, 26, 40, 46, 22, 45\}$. Students who share courses are connected, and the walk explores the local neighborhood structure.

3.4 Learning Student Embeddings

The generated random walks are used to learn low-dimensional student representations using a skip-gram objective. Let $z_v \in \mathbb{R}^d$ denote the embedding of student v , where d is the embedding dimension. For a student v , the objective encourages students appearing within a context window of v to have similar representations. Figure 3 illustrates the CBOW and Skip-gram architectures used in this learning stage.

Given a context window of size r , the optimization objective can be expressed as

$$\max_{\{z_v\}} \sum_{v \in V} \sum_{u \in C(v)} \log P(u | v),$$

where $C(v)$ denotes the set of students occurring within the specified context window of v in the generated random walks.

The conditional probability can be modeled using the softmax function:

$$P(u | v) = \frac{\exp(z_u^\top z_v)}{\sum_{x \in V} \exp(z_x^\top z_v)}.$$




Fig. 3: The CBOW architecture predicts the current word based on the context, and the Skip-gram predicts surrounding words given the current word [16]. In DeepWalk, nodes play the role of words and random-walk sequences play the role of sentences.

In practice, negative sampling can be used to make the optimization computationally efficient. The resulting embedding matrix

$$Z = \begin{bmatrix} z_1^\top \\ z_2^\top \\ \vdots \\ z_n^\top \end{bmatrix} \in \mathbb{R}^{n \times d}$$

provides a continuous representation of all students.

The important distinction from the original enrollment representation is that each vector z_i is learned from the structural context of student s_i in the entire co-enrollment network. Consequently, the representation can capture indirect and higher-order similarities between students.

Algorithm 2 summarizes the overall procedure: biased random walks are generated from the co-enrollment graph (line 1), and the Skip-Gram model is trained on these walks to produce the student embedding matrix Φ (line 2 of algorithm 2).

3.5 Student Sectioning by Clustering

Once the student embeddings have been learned, student sectioning is formulated as a clustering problem. Let k denote the desired number of sections. A clustering algorithm partitions the set of student embeddings

$$Z = \{z_1, z_2, \dots, z_n\}$$

into k clusters,

Algorithm 1 GenerateRandomWalks

Require: Graph G , walk length t , number of walks per node γ

Ensure: A list $\mathcal{W}$ containing all sequences of random walks

```

1:  $\mathcal{W} \leftarrow []$  ▷ Initialize list to store walks
2: for each graph node as  $start$  do
3:   for  $k = 1$  to  $\gamma$  do ▷ Number of walks per node
4:      $walk \leftarrow [start]$  ▷ Start the walk at the
       current node
5:     for  $j = 1$  to  $t$  do ▷ Walk length
6:        $start \leftarrow$  randomly chosen node from list of
       neighbors of  $start$ 
7:        $walk \leftarrow walk + [start]$  ▷ Append next
       node to walk
8:     end for
9:      $\mathcal{W} \leftarrow \mathcal{W} + [walk]$  ▷ Store the walk
10:  end for
11: end for
12: return  $\mathcal{W}$ 
```

Algorithm 2 Student Representation Generation using Node2Vec

Require: graph $G(V, E)$ constructed from student-course matrix, where students are its nodes. The input parameters are:

walk length t
number of walks per node γ
embedding size d
window size w
epochs ϵ

Ensure: matrix of vertex (student) representations $\Phi \in \mathbb{R}^{|V| \times d}$

```

1:  $\mathcal{W} \leftarrow \text{GenerateRandomWalks}(G, t, \gamma)$  ▷ Algorithm 1
2:  $\Phi \leftarrow \text{TrainWord2Vec}(\mathcal{W}, d, w, \epsilon)$  ▷ Sect. 3.4
3: return  $\Phi$ 
```

$$\mathcal{S} = \{S_1, S_2, \dots, S_k\},$$

where each cluster S_j corresponds to a student section.

The proposed framework is deliberately independent of a particular clustering algorithm. This allows us to investigate whether the improvement comes from the

learned graph representations rather than from a specific clustering technique. In our experiments, we therefore evaluate several widely used clustering methods, including K-means, affinity propagation, Gaussian mixture models, and hierarchical clustering.

For K-means, for example, the clustering objective is

$$\min_{\{S_j\}_{j=1}^k} \sum_{j=1}^k \sum_{z_i \in S_j} \|z_i - \mu_j\|_2^2,$$

where μ_j denotes the centroid of section S_j .

The resulting clusters define the student sections used in the subsequent evaluation.

3.6 Baseline Representation

To determine whether graph representation learning provides an advantage over conventional student representations, we establish a baseline based directly on the student-course enrollment matrix.

For each student s_i , the corresponding row of the enrollment matrix,

$$x_i = (X_{i1}, X_{i2}, \dots, X_{im}),$$

is used as the student representation. Thus, students are represented directly by their binary course enrollment patterns without constructing a graph or learning an embedding.

The same clustering algorithms are then applied to these enrollment vectors. This provides a controlled comparison because the clustering stage remains unchanged and the principal difference between the two approaches is the representation used to describe students.

Accordingly, the comparison considers two alternative pipelines:

Enrollment representation $\rightarrow$ Clustering,

and

Graph representation $\rightarrow$ Embedding $\rightarrow$ Clustering.

This design allows the contribution of graph representation learning to be evaluated independently of the choice of clustering algorithm.

3.7 Hyperparameter Configuration

The graph embedding procedure contains several hyperparameters that may affect the quality of the learned representations. The principal parameters considered in this study include the embedding dimension d , the random-walk length L , the number of walks generated per student R_v , and the context-window size r .

The embedding dimension controls the dimensionality of the learned student representations. A larger value provides greater representational capacity but may also increase computational cost and the risk of overfitting. The walk length determines how far the random walks explore the co-enrollment network, while the number of walks controls the amount of structural context available during training. The context-window size determines the range of neighboring positions considered when learning the embedding.

Rather than relying on a single arbitrary configuration, we conduct sensitivity experiments over these parameters. The purpose is not to optimize the method specifically for a particular dataset, but to examine whether the observed improvement is robust across reasonable parameter settings.

3.8 Evaluation of Student Sectioning

The quality of the resulting student sections is evaluated using clustering-based internal validity measures. Since the primary objective is to create coherent groups of students, the Silhouette Score is used as the main evaluation metric.

For a student i , let $a(i)$ denote the average distance between i and the other students in the same section, and let $b(i)$ denote the minimum average distance between i and the students belonging to another section. The silhouette coefficient of student i is defined as

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}.$$

The overall Silhouette Score is obtained by averaging this coefficient over all students:

$$S = \frac{1}{n} \sum_{i=1}^n s(i).$$

The Silhouette Score ranges from -1 to 1 , with larger values indicating better within-section cohesion and between-section separation.

To complement the primary evaluation, we also examine clustering stability across multiple random seeds. In particular, the Adjusted Rand Index (ARI) is used to quantify the agreement between clustering solutions obtained under different random initializations. High ARI values indicate that the learned representation leads to stable section assignments rather than highly variable partitions.

Finally, statistical significance tests are employed to determine whether the observed differences between the proposed graph-based representation and the baseline enrollment representation are unlikely to be attributable to random variation. Effect-size measures are also reported to quantify the practical magnitude of the observed improvements.

4 Experimental Results

We evaluate the proposed graph representation learning framework for student sectioning using a real-world university dataset and a series of experiments designed to examine both its effectiveness and robustness. The experiments address four main questions: (i) whether graph-based representations improve student clustering compared with the conventional student-course representation, (ii) whether the improvement can be explained solely by dimensionality reduction, (iii) whether direct graph clustering provides comparable benefits, and (iv) how sensitive the proposed approach is to its main hyperparameters and stochastic components.

4.1 Dataset and Experimental Setup

We use the real-world student-course enrollment dataset introduced in [2]. To evaluate the student sectioning problem, we select six relatively large courses containing 74, 51, 48, 49, 52, and 67 students, respectively. These courses provide a range of realistic sectioning scenarios while retaining sufficient enrollment information to construct meaningful student co-enrollment graphs.

The proposed method constructs a weighted student co-enrollment graph from the preprocessed enrollment data and learns student representations using Node2Vec

random walks and the Skip-Gram model. The resulting embeddings are subsequently provided to a clustering algorithm for student sectioning. Unless otherwise stated, KMeans is used as the downstream clustering algorithm.

As the conventional baseline, we use the original binary student-course enrollment representation (after constant-feature removal), referred to as the *Student-Course Representation* or *BoW* representation. To examine whether the observed improvement is primarily attributable to dimensionality reduction, we additionally evaluate a PCA-based baseline in which the conventional representation is projected to the same two-dimensional space used by the graph embedding methods before applying KMeans. We also evaluate spectral clustering directly on the student co-enrollment graph as a graph-based baseline that does not use an intermediate learned embedding.

Unless otherwise stated, the default Node2Vec parameters are $p = 1$, $q = 1$ (unbiased random walk), $t = 10$, $\gamma = 80$, $d = 2$, $w = 5$, and $\epsilon = 30$, where p and q are the Node2Vec bias parameters, t denotes walk length, γ the number of walks generated per node, d the embedding dimension, w the Skip-Gram context-window size, and ϵ the number of Word2Vec training epochs. A parameter sensitivity experiment over $p, q \in \{0.5, 1.0, 2.0\}$ is also conducted. Experiments involving stochastic components are repeated over 20 random seeds. The source code and datasets used in this paper are available as supplementary material.

The primary evaluation metric is the Silhouette Score [17], which considers both intra-cluster cohesion and inter-cluster separation. Higher values indicate better-defined clusters. We additionally report the Davies–Bouldin Index (DBI), for which lower values are better, and the Calinski–Harabasz (CH) Index, for which higher values are better.

4.2 Overall Clustering Performance and Baseline Comparison

We compare Node2Vec with the conventional student-course representation, PCA+KMeans, spectral clustering, and DeepWalk (which corresponds to Node2Vec with $p = 1$, $q = 1$) to determine whether the observed improvement can be attributed simply to dimensionality reduction, to the specific random-walk strategy, or to direct exploitation of graph structure.

Table 1: Silhouette Score ($\uparrow$ higher is better) comparison across five methods for student sectioning with the revised graph construction. PCA+KMeans applies PCA to reduce the Student-Course representation to $d = 2$ before KMeans clustering. Spectral applies spectral clustering directly on the co-enrollment graph. DeepWalk uses the learned $d = 2$ graph embeddings from unbiased random walks followed by KMeans (equivalent to Node2Vec with $p=1, q=1$). Node2Vec uses $p=1, q=1$ unless otherwise noted.

Course	BoW + KMeans	PCA + KMeans	Spectral	DeepWalk	Node2Vec
1	0.099	0.357	0.154	0.570	0.535
2	0.231	0.603	0.340	0.617	0.599
3	0.142	0.425	0.148	0.519	0.576
4	0.128	0.428	0.181	0.574	0.624
5	0.116	0.425	0.255	0.766	0.582
6	0.202	0.524	0.199	0.633	0.611
Average	0.153	0.460	0.213	0.613	0.611




Fig. 4: Silhouette scores ($\uparrow$ higher is better) for five student sectioning methods across six courses: BoW + KMeans, PCA + KMeans, Spectral Clustering, DeepWalk (unbiased random walk), and Node2Vec ($p=1, q=1$).

Figure 4 and Table 1 summarize the Silhouette scores obtained by the five approaches across the six selected courses.

Both graph embedding methods achieve substantially higher average Silhouette Scores than the conventional BoW representation, spectral clustering, and (in the case of DeepWalk) also higher than PCA+KMeans. DeepWalk

achieves the highest average Silhouette Score of 0.613, compared with 0.611 for Node2Vec ($p=1, q=1$), 0.460 for PCA+KMeans, 0.213 for spectral clustering, and 0.153 for the conventional BoW representation.

Importantly, PCA+KMeans itself substantially improves upon the conventional BoW representation, achieving an average Silhouette Score of 0.460. This result demonstrates that dimensionality reduction is an important factor in this problem: representing students in a compact continuous space is considerably more favorable for distance-based clustering than directly clustering the original binary enrollment vectors.

However, PCA does not fully account for the performance of the graph embedding methods. DeepWalk improves the average Silhouette Score from 0.460 with PCA+KMeans to 0.613, corresponding to a relative improvement of approximately 33.3%. Node2Vec ($p=1, q=1$) achieves 0.611, a relative improvement of 32.8% over PCA+KMeans. Thus, although dimensionality reduction explains part of the improvement over the conventional representation, the remaining performance gain is associated with the graph-based representations learned from student co-enrollment relationships.

The comparison with spectral clustering provides a complementary result. Spectral clustering achieves an average Silhouette Score of 0.213, which is only modestly better than the conventional BoW+KMeans baseline (0.153). Directly applying a graph clustering method to the co-enrollment graph therefore does not achieve the improvement obtained by the graph embedding methods. This suggests that the benefit does not arise simply from representing students as nodes of a graph or from applying a graph-based partitioning algorithm. Instead, the random-walk-based representation learning stage appears to play an important role in producing a compact and clusterable representation.

An important observation is that Node2Vec with $p=1$ and $q=1$ is theoretically equivalent to an unbiased random walk, which corresponds to the DeepWalk setting. The near-identical average Silhouette Scores of DeepWalk (0.613) and Node2Vec ($p=1, q=1$) (0.611) are consistent with this theoretical equivalence. The small per-course differences are attributable to different random-number streams in the respective implementations. The sensitivity experiment (Section 4.3) further shows that adjusting

Table 2: Student co-enrollment graph statistics after constant-feature removal. n is the number of students, m is the number of remaining courses, $|E|$ is the number of edges, and $\bar{d}$ is the average weighted degree.

Course	n	m	$ E $	Density	Components
1	74	34	1,828	0.677	3
2	51	33	866	0.679	1
3	48	27	894	0.793	2
4	49	27	1,043	0.887	1
5	52	31	912	0.688	1
6	67	32	1,325	0.599	1
Average	56.8	30.7	1,145	0.721	1.7

the Node2Vec parameters to $p=1$, $q=0.5$ yields a modest improvement (average Silhouette 0.622), suggesting that a slight BFS-like walk bias can marginally improve embedding quality on these graphs.

Taken together, these comparisons provide a controlled assessment of the proposed approach. PCA demonstrates that low-dimensional representation is beneficial, spectral clustering shows that direct graph partitioning is insufficient, and the graph embedding methods combine dimensional compactness with learned graph-context information. The results therefore support the interpretation that the performance gain is not solely a consequence of dimensionality reduction, but also reflects information captured from the relational structure of the student co-enrollment graph.

Table 2 summarizes the graph structure after constant-feature removal.

Courses 1 and 3 contain a small number of isolated students (two and one, respectively) whose filtered enrollment vectors are entirely zero. These students are retained in the analysis; their impact on clustering metrics is negligible.

4.3 Node2Vec Parameter Sensitivity

We conduct a parameter sensitivity experiment over $p \in \{0.5, 1.0, 2.0\}$ and $q \in \{0.5, 1.0, 2.0\}$, keeping all other parameters at their default values ($t=10$, $\gamma=80$, $d=2$, $w=5$, $\epsilon=30$). Table 3 reports the average Silhouette Score across the six courses for each configuration.

Table 3: Node2Vec parameter sensitivity: average Silhouette Score across six courses for different (p, q) configurations.

q	$p = 0.5$			$p = 1.0$			$p = 2.0$	
	0.5	1.0	2.0	0.5	1.0	2.0	0.5	1.0
Silhouette	0.620	0.616	0.614	0.622	0.611	0.605	0.615	0.611

Table 4: Effect of embedding dimension d on clustering quality, averaged across the six courses. Silhouette score ($\uparrow$ higher is better), Davies–Bouldin Index ($\downarrow$ lower is better), and Calinski–Harabasz Index ($\uparrow$ higher is better).

Dimension (d)	Silhouette ($\uparrow$)	DBI ($\downarrow$)	CH ($\uparrow$)
1	0.552	0.593	95.2
2	0.579	0.564	106.6
3	0.377	1.011	40.0
5	0.204	1.707	14.9
10	0.088	2.799	6.3

The best configuration is $p=1.0$, $q=0.5$ with an average Silhouette Score of 0.622. The neutral configuration $p=1$, $q=1$ (equivalent to unbiased random walks) achieves 0.611. The sensitivity range is modest (0.605–0.622), and lower values of q (BFS-like exploration) tend to perform slightly better. The parameter p has a smaller effect than q within the tested range. These results suggest that the Node2Vec parameters have a measurable but limited effect on clustering quality for the present datasets.

4.4 Embedding Dimension and Hyperparameter Sensitivity

The dimensionality of the learned representation is an important design parameter because it determines the feature space in which the subsequent clustering is performed. We therefore examine embedding dimensions $d \in \{1, 2, 3, 5, 10\}$. Figure 5 shows the resulting Silhouette scores, while Table 4 reports three complementary clustering metrics.

The results consistently identify $d = 2$ as the best-performing configuration among the tested embedding




Fig. 5: Silhouette scores ($\uparrow$ higher is better) for the graph embedding method across different embedding dimensions. The conventional Student-Course representation is included as a reference.

dimensions. It achieves the highest Silhouette Score (0.579), the lowest DBI (0.564), and the highest CH Index (106.6). The one-dimensional representation also performs relatively well, whereas the clustering quality decreases substantially for dimensions greater than two. At $d = 10$, the Silhouette Score falls to 0.088.

These results indicate that, for the present student sectioning dataset, a compact two-dimensional representation provides the most effective balance between the information captured by the graph embedding and the requirements of the subsequent clustering task. This finding should be interpreted as an empirical property of the present experimental setting rather than as a general claim that student relationships are intrinsically two-dimensional. The appropriate embedding dimensionality may depend on the size and structural characteristics of the underlying enrollment graph.

We next examine the sensitivity of the graph embedding method to its principal random-walk and Skip-Gram parameters. The effects of walk length t , number of walks per node γ , and context-window size w are evaluated independently, while the remaining parameters are fixed at their default values.

Walk length has a noticeable effect on clustering quality, as shown in Table 5. The highest Silhouette Score is obtained with $t = 5$ (0.609), followed by $t = 10$ (0.579), with performance gradually decreasing as the walks become longer. This result suggests that relatively

Table 5: Sensitivity to walk length t ($\uparrow$ higher Silhouette is better; averaged across six courses, $d = 2$, $\gamma = 80$, $w = 5$).

Walk Length (t)	5	10	20	40	80
Silhouette ($\uparrow$)	0.609	0.579	0.574	0.565	0.538

Table 6: Sensitivity to the number of walks per node γ ($\uparrow$ higher Silhouette is better; averaged across six courses, $d = 2$, $t = 10$, $w = 5$).

Walks (γ)	10	20	40	80	160
Silhouette ($\uparrow$)	0.567	0.569	0.571	0.579	0.562

Table 7: Sensitivity to context-window size w ($\uparrow$ higher Silhouette is better; averaged across six courses, $d = 2$, $t = 10$, $\gamma = 80$).

Window (w)	1	2	5	10	20
Silhouette ($\uparrow$)	0.555	0.557	0.579	0.584	0.581

local co-enrollment relationships are particularly informative for the present sectioning task. Longer walks introduce increasingly distant graph relationships into the training corpus, which may reduce the relevance of the resulting context for student grouping.

The number of walks per node has a smaller effect than embedding dimension and walk length, as shown in Table 6. Performance increases gradually from $\gamma = 10$ to $\gamma = 80$, where the highest score (0.579) is obtained, and then decreases slightly at $\gamma = 160$. The results therefore suggest that $\gamma = 80$ provides a reasonable trade-off between sampling density and computational cost for the present dataset.

The context-window size has the smallest effect among the examined parameters, as shown in Table 7. The Silhouette Score remains within the relatively narrow range of 0.555–0.584, with the best result obtained at $w = 10$. Thus, the performance of the graph embedding method is relatively insensitive to the precise choice of the context-window size within the tested range.




Fig. 6: Silhouette score ($\uparrow$ higher is better) comparison across the six courses for different clustering algorithms applied to the conventional Student-Course representation and Node2Vec embeddings.

Overall, the sensitivity experiments show that embedding dimensionality and walk length have the largest effects on clustering quality for DeepWalk, whereas the number of walks and context-window size have comparatively smaller effects. A similar pattern is expected for Node2Vec, as the two methods share the same embedding procedure and differ only in the walk-generation strategy. The Silhouette Score varies from 0.088 to 0.579 across the tested embedding dimensions, compared with a range of 0.071 across walk lengths and only 0.029 across context-window sizes. For consistency with the main baseline comparison, we retain $d = 2$, $t = 10$, $\gamma = 80$, and $w = 5$ as the default configuration. The sensitivity analysis nevertheless indicates that a shorter walk length, particularly $t = 5$, can provide further improvement for this dataset.

4.5 Clustering Algorithms and Representation Quality

To determine whether the observed improvement depends on a particular clustering algorithm, we evaluate the learned Node2Vec representations using four clustering methods: KMeans, affinity propagation, Gaussian mixture models (GMM), and hierarchical clustering. Figure 6 summarizes the resulting Silhouette scores across the six courses.

Node2Vec generally produces higher clustering quality across the evaluated clustering algorithms. Although the magnitude of the improvement varies between algorithms and courses, the advantage is not restricted to KMeans. This provides evidence that the observed improvement is

primarily associated with the learned student representation rather than with a particular interaction between Node2Vec and a specific clustering algorithm.

This result is consistent with the formulation of the proposed framework, in which graph representation learning is used as a preprocessing stage and the downstream clustering procedure is deliberately kept unchanged. The purpose is therefore not to introduce a specialized clustering algorithm, but to provide a representation in which students with structurally similar enrollment patterns become more suitable for subsequent grouping.

Combined with the PCA comparison reported in Table 1, these results provide a more complete interpretation of the role of representation quality. PCA substantially improves clustering relative to the original high-dimensional enrollment representation, showing that dimensionality reduction is beneficial. Nevertheless, Node2Vec achieves a higher average Silhouette Score than PCA+KMeans. The additional improvement suggests that the graph-based embedding captures relational information that is not fully recovered by applying a generic linear dimensionality-reduction technique to the enrollment matrix.

4.6 Clustering Stability and Additional Evaluation Metrics

In addition to clustering quality, reproducibility is important in student sectioning because unstable assignments may lead to inconsistent student groups across repeated runs. We therefore evaluate clustering stability across 20 random seeds using the Adjusted Rand Index (ARI). The results are presented in Table 8.

The Node2Vec representation produces highly stable clusterings across repeated runs. KMeans achieves an average ARI of 0.986, while agglomerative clustering achieves an ARI of 1.000 across the six courses. GMM exhibits somewhat greater variability, with an average ARI of 0.879, but remains relatively stable. In comparison, KMeans applied to the conventional Student-Course Representation has a lower average ARI of 0.856, with particularly low stability for Course 1 (0.426).

We further evaluate clustering quality using the Davies–Bouldin Index (DBI) and the Calinski–Harabasz (CH) Index. Figures 7 and 8 present the corresponding comparisons.

Table 8: Clustering stability measured by Adjusted Rand Index (ARI) across 20 random seeds ($\uparrow$ higher is better).

Course	Node2Vec , ARI			BoW
	KMeans	GMM	Agglom.	KMeans
1	1.000	1.000	1.000	0.426
2	1.000	0.845	1.000	1.000
3	1.000	0.760	1.000	0.854
4	0.917	1.000	1.000	0.992
5	1.000	1.000	1.000	0.865
6	1.000	0.670	1.000	1.000
Average	0.986	0.879	1.000	0.856




Fig. 7: Davies-Bouldin Index ($\downarrow$ lower is better) across the six courses for different clustering methods and representations.

The DBI and CH results provide complementary evidence regarding the structure of the obtained clusters. Across the evaluated courses and clustering methods, Node2Vec generally produces more favorable DBI and CH values than the conventional representation. These results are consistent with the Silhouette analysis and indicate that the observed improvement is reflected in multiple internal clustering criteria rather than in a single evaluation metric.

4.7 Statistical Significance

We assess whether the observed improvements of the graph embedding methods over the baselines are statistically significant. Wilcoxon signed-rank tests are performed using paired per-course means obtained from 20 random seeds. Because there are only six courses,




Fig. 8: Calinski-Harabasz Index ($\uparrow$ higher is better) across the six courses for different clustering methods and representations.

the independent sample size is $n = 6$. This limits the statistical power of the tests.

Node2Vec versus conventional representation.

The Wilcoxon signed-rank test comparing Node2Vec ($p=1, q=1$) against the conventional BoW representation yields a raw p -value of 0.031 (the minimum achievable for a two-sided test with $n = 6$). The effect size is $r = 0.899$ (large), and Cliff's $\delta = 1.0$, meaning that Node2Vec outperforms the conventional representation on all six courses. However, after Holm correction for multiple comparisons, the adjusted p -value is 0.188, which does not reach significance at $\alpha = 0.05$. The large effect sizes nevertheless indicate a consistent and practically meaningful improvement.

Node2Vec versus PCA+KMeans.

The comparison of Node2Vec ($p=1, q=1$) against PCA+KMeans also yields a raw p -value of 0.031, with the same effect sizes ($r = 0.899, \delta = 1.0$). Node2Vec outperforms PCA+KMeans on five of the six courses (PCA+KMeans achieves a higher score on Course 2). After Holm correction, the adjusted p -value is 0.188. These results indicate that the graph-based representation provides a consistent advantage over linear dimensionality reduction, but the small number of datasets prevents the difference from reaching formal statistical significance after correction.

Node2Vec versus DeepWalk.

Node2Vec ($p=1$, $q=1$) and DeepWalk achieve near-identical average Silhouette Scores (0.611 vs. 0.613). The Wilcoxon test yields $p = 1.0$, confirming that there is no meaningful difference between the two methods at this configuration. This is consistent with the theoretical expectation that Node2Vec with $p=1$, $q=1$ performs unbiased random walks equivalent to DeepWalk.

In summary, the effect sizes for graph methods versus baselines are large (Cliff’s $\delta = 1.0$), but the formal statistical significance does not survive Holm correction due to the small number of independent datasets ($n = 6$). We recommend interpreting the results primarily through effect sizes and the consistency of improvement across courses rather than relying solely on p -values.

4.8 Embedding Visualization


To provide a qualitative interpretation of the quantitative results, we visualize the learned representations for a sample course. The visualizations compare the conventional Student-Course representation with the two-dimensional Node2Vec representation. For the conventional Student-Course representation, which is defined in a higher-dimensional space, t-Distributed Stochastic Neighbor Embedding (t-SNE) is used solely for visualization. In contrast, Node2Vec directly produces a two-dimensional embedding with $d = 2$, which is the representation used for clustering in the corresponding experiments. Therefore, no additional dimensionality-reduction step is applied to the Node2Vec embeddings.

Figure 9 presents the representations obtained for Course 5. The visualization of the conventional representation shows a relatively overlapped and dispersed structure, whereas the Node2Vec representation exhibits more compact and visually separated groups. This qualitative difference is consistent with the substantially higher Silhouette Score obtained by Node2Vec for this course.


These visualizations provide qualitative support for the quantitative results, while the distinction between the two visualization procedures should be noted: t-SNE is applied only to the high-dimensional conventional representation, whereas the Node2Vec points are plotted directly in their learned two-dimensional embedding space.

4.9 Runtime Analysis

Finally, we examine the computational cost of the complete Node2Vec pipeline. Table 9 reports the runtime of graph construction, random-walk generation, Word2Vec training, and the complete pipeline for each course. The complete pipeline requires less than nine seconds for all six courses. Word2Vec training is the dominant component of the runtime, accounting for approximately 80% of the total execution time, whereas graph construction and walk generation require comparatively little time.



(a) t-SNE of the *Student-Course Rep.*.



(b) Node2Vec embeddings ($d = 2$).

Fig. 9: Visualization of student representations for Course 5. (a) The conventional Student-Course representation is defined in a 32-dimensional space and is projected to two dimensions using t-SNE solely for visualization. (b) Node2Vec directly produces a two-dimensional representation, so no additional dimensionality-reduction method is required. In both panels, students are displayed according to their resulting section assignments.

Table 9: Runtime analysis in seconds for the complete Node2Vec pipeline using $d = 2$, $t = 10$, $\gamma = 80$, $w = 5$, and $\epsilon = 30$.

Course	Students	Graph	Walks	Word2Vec	Total
1	74	0.024	0.263	8.067	8.355
2	51	0.011	0.168	5.471	5.650
3	48	0.009	0.144	4.584	4.738
4	49	0.070	0.145	4.785	5.000
5	52	0.009	0.153	4.830	4.992
6	67	0.016	0.218	6.564	6.798

For a dataset with n students and m courses, construction of the student co-enrollment graph from the enrollment matrix requires comparing all student pairs over all courses, yielding a cost of $O(n^2m)$. Random-walk generation requires $O(n\gamma t)$ operations, where γ is the number of walks per node and t is the walk length. The cost of Skip-Gram training is proportional to the total number of walk-context pairs and the embedding dimension, approximately $O(n\gamma td)$, where d is the embedding dimension. For the relatively small to moderate student populations considered in this study, the overall computational cost remains modest.

Overall, the experiments demonstrate that the Node2Vec representation provides substantially improved clustering quality relative to the conventional Student-Course representation in the evaluated setting. The improvement is observed across multiple clustering algorithms and evaluation metrics, while the stability analysis indicates that the resulting clusterings are highly reproducible. The sensitivity experiments further show that effective representations can be obtained in a very low-dimensional space. The runtime analysis indicates that the complete pipeline can be executed within a few seconds for the course sizes considered in this study.

5 Discussion

The experimental results provide consistent evidence that graph representation learning can improve the quality of student sectioning. Across the six evaluated courses, Node2Vec produces markedly higher clustering quality than the conventional student-course representation and

spectral clustering. The improvement remains evident across different evaluation metrics (Silhouette, DBI, CH) and is highly stable across 20 random seeds. The results therefore suggest that the representation of student relationships plays an important role in the effectiveness of subsequent clustering.

A key explanation for this improvement is the difference between the two representations. The conventional student-course representation describes each student through direct course enrollments. Although this representation contains useful information, relationships between students are represented only implicitly through shared binary features. In contrast, the co-enrollment graph explicitly models students as nodes and their shared course enrollments as weighted edges. Random walks on this graph expose the embedding model to sequences of structurally related students, allowing the resulting representations to encode higher-order proximity and contextual relationships that are not directly expressed in the original enrollment matrix. The substantially higher Silhouette scores observed with Node2Vec are consistent with this interpretation.

The additional baseline experiments provide an important control for this interpretation. Applying PCA to the conventional student-course representation and reducing it to the same two-dimensional space improves the average Silhouette Score from 0.153 to 0.460. This demonstrates that dimensionality reduction itself contributes substantially to clustering quality and therefore represents an important factor in the comparison. However, Node2Vec ($p=1$, $q=1$) achieves an average score of 0.611, indicating that its advantage cannot be explained solely by reducing the dimensionality of the original enrollment matrix. Similarly, spectral clustering applied directly to the student co-enrollment graph achieves an average Silhouette Score of 0.213, which is only modestly better than the conventional representation and substantially below the graph embedding methods. Within the present experimental setting, these comparisons support the interpretation that learning a task-oriented low-dimensional representation from the graph structure is more effective than either dimensionality reduction of the original matrix or direct spectral clustering of the graph.

An important observation is that Node2Vec with $p=1, q=1$ is theoretically equivalent to the unbiased random walk used by DeepWalk. The near-identical average Silhouette Scores of DeepWalk (0.613) and Node2Vec (0.611) confirm this equivalence empirically. The sensitivity experiment shows that the best tested configuration is $p=1, q=0.5$ (0.622), suggesting that a modest BFS-like walk bias can slightly improve embedding quality. The overall parameter sensitivity is limited, with Silhouette Scores ranging from 0.605 to 0.622 across the nine tested (p, q) configurations.

Another important finding is that the best performance is obtained with a very low-dimensional embedding, particularly $d = 2$. This result should not be interpreted as evidence that low-dimensional embeddings are universally preferable for graph representation learning. Rather, it reflects the characteristics of the present student sectioning problem and dataset. The task considered here involves relatively small student groups and binary sectioning. Under these conditions, a compact representation appears sufficient to preserve the dominant structural patterns relevant to clustering. The two-dimensional representation also provides a practical advantage because the learned embeddings can be directly visualized without requiring an additional dimensionality-reduction method.

The parameter sensitivity experiments provide further insight into the behavior of the proposed framework. Shorter random walks generally perform better than longer walks, indicating that relatively local co-enrollment relationships are particularly informative for the investigated sectioning problem. The method also shows comparatively limited sensitivity to the number of walks and context-window size over the investigated ranges. These observations suggest that the method can be applied without extensive hyperparameter tuning to datasets with characteristics similar to those studied here.

The improvement is also largely independent of the clustering algorithm. Node2Vec consistently provides better clustering quality than the conventional representation when combined with KMeans, affinity propagation, Gaussian mixture models, and hierarchical clustering. This observation is important because the proposed framework does not depend on introducing a specialized clustering procedure. Instead, graph representation learning can

serve as a preprocessing stage, after which an institution may select a clustering algorithm according to its operational requirements.

The stability analysis provides complementary evidence. Node2Vec produces highly consistent cluster assignments across 20 random seeds, with a mean course-level standard deviation of 0.013. Such stability is desirable in student sectioning because substantial changes in assignments caused solely by stochastic initialization could complicate administrative planning. For comparison, DeepWalk shows a slightly higher variability (mean standard deviation 0.020).

The results should nevertheless be interpreted within the scope of the experimental setting. The evaluation focuses on internal clustering criteria rather than directly optimizing or measuring downstream timetabling objectives. A higher Silhouette Score indicates more compact and better-separated clusters in the representation space, but it does not by itself guarantee fewer timetable conflicts, better room utilization, or improved feasibility of the complete scheduling problem. Furthermore, although the effect sizes are large (Cliff’s $\delta = 1.0$), the improvements do not reach statistical significance after Holm correction due to the small number of independent datasets ($n = 6$). The present study therefore establishes the effectiveness of graph-based representation learning as a clustering-based approach to student sectioning, while the direct impact of the resulting sections on complete timetabling remains an important direction for future investigation.

6 Conclusion

This paper revisits the student sectioning problem from a graph representation learning perspective. Instead of representing students solely through a binary student-course matrix, we construct a weighted student co-enrollment graph and use Node2Vec to obtain low-dimensional student embeddings through biased random walks. These embeddings are subsequently used as inputs to standard clustering algorithms for section formation.

The experimental results on six real-world university courses demonstrate that Node2Vec improves clustering quality compared with the conventional student-course representation and spectral clustering. With the neutral

configuration ($p=1$, $q=1$), Node2Vec achieves an average Silhouette Score of 0.611, compared with 0.153 for the conventional representation, 0.460 for PCA followed by KMeans, and 0.213 for spectral clustering. The best tested configuration ($p=1$, $q=0.5$) achieves 0.622. The improvement is consistent across all six courses, with large effect sizes (Cliff’s $\delta = 1.0$). Additional evaluation using the Davies–Bouldin and Calinski–Harabasz indices provides complementary evidence that the graph-based representations produce more compact and better-separated clusters.

The baseline experiments clarify the source of the improvement. PCA applied to the conventional student-course representation and reduced to $d = 2$ achieves an average Silhouette Score of 0.460, demonstrating that dimensionality reduction alone substantially improves the baseline. However, Node2Vec reaches 0.611, indicating that its performance cannot be attributed solely to dimensionality reduction. Spectral clustering applied directly to the co-enrollment graph obtains an average score of 0.213. These comparisons suggest that, for the datasets and experimental protocol considered here, learning a low-dimensional representation from the graph structure provides a more effective basis for student clustering than either direct clustering of the original representation or direct spectral graph clustering.

The experiments also provide useful insights into the behavior of the framework. A two-dimensional embedding provides the best performance among the evaluated embedding dimensions. The sensitivity analysis shows that the Node2Vec parameters have a measurable but modest effect, with Silhouette Scores ranging from 0.605 to 0.622 across the tested configurations. Relatively short random walks are preferable for the investigated student co-enrollment graphs. The stability experiments show that the resulting cluster assignments are highly reproducible across 20 random seeds (mean standard deviation of 0.013).

From a practical perspective, the proposed framework is computationally lightweight for the problem sizes considered in this study. The complete pipeline requires a few seconds for each of the evaluated courses under the default configuration. Node2Vec has comparable runtime to an unbiased random walk because the biased transition probability calculation is negligible relative to the

Word2Vec training time. This indicates that graph representation learning can be incorporated into student sectioning workflows without substantial computational overhead for datasets of comparable size.

Several directions naturally follow from this work. First, the current study could be extended to sparser co-enrollment graphs where the Node2Vec q parameter would have a greater effect on walk behavior. Second, graph neural network methods such as GraphSAGE could be investigated to incorporate additional student attributes alongside graph structure. Third, dynamic graph construction could model changes in enrollment during registration and add/drop periods. Fourth, integrating graph-based sectioning with constraint programming or other timetabling optimization techniques could allow the impact of improved student grouping to be evaluated directly in terms of timetable quality. Finally, validation on larger and multi-institutional datasets would provide stronger evidence regarding the generalizability of the approach.

In summary, the results indicate that the student co-enrollment graph contains useful relational information that is only partially captured by conventional student-course representations. Exploiting this information through Node2Vec provides a simple and computationally efficient way to obtain more coherent student groupings. Node2Vec outperforms conventional representations and PCA-based dimensionality reduction on all six courses. Although the improvements do not survive Holm correction for multiple comparisons (due to the small number of independent datasets), the large effect sizes and consistent improvement across courses suggest that graph-based representation learning is a promising approach to student sectioning that warrants further investigation with larger datasets.

Competing interests

The author has no relevant financial or non-financial interests to disclose.

References

- [1] Amintoosi, M., Yazdi, H.S., Haddadnia, J.: Fuzzy student sectioning. In: PATAT04: Practice and Theory of Automated Timetabling, USA, pp. 421–424

- (2004)
- [2] Amintoosi, M., Haddadnia, J.: Feature selection in a fuzzy student sectioning algorithm. In: Burke, E., Trick, M. (eds.) *Practice and Theory of Automated Timetabling V*, pp. 147–160. Springer, Berlin, Heidelberg (2005). *Lecture Notes in Computer Science (LNTCS, volume 3616)*
 - [3] Amintoosi, M., Fathy, M., Mozayani, N., Rahmani, A.T.: A fish school clustering algorithm: Applied to student sectioning problem. *Dynamics of Continuous Discrete & Impulse Systems, series B: Applications and Algorithms* **2**, 696–699 (2007). *Post Proceeding of LSMS2007, Life System Modeling and Simulation 2007, China*
 - [4] Aubin, J., Ferland, J.A.: A large scale timetabling problem. *Computers & Operations Research* **16**(1), 67–77 (1989)
 - [5] Hertz, A., Robert, V.: Constructing a course schedule by solving a series of assignment type problems. *European Journal of Operational Research* **108**(3), 585–603 (1998)
 - [6] Carter, M.W.: A comprehensive course timetabling and student scheduling system at the university of waterloo. In: Burke, E., Erben, W. (eds.) *Practice and Theory of Automated Timetabling III*, pp. 64–82. Springer, Berlin, Heidelberg (2001)
 - [7] Alvarez-Valdes, R., Crespo, E., Tamarit, J.M.: Assigning students to course sections using tabu search. *Annals of Operations Research* **96**(1), 1–16 (2000)
 - [8] Müller, T., Murray, K.: Comprehensive approach to student sectioning. *Annals of Operations Research* **181**(1), 249–269 (2010) <https://doi.org/10.1007/s10479-010-0735-9>
 - [9] Perozzi, B., Al-Rfou, R., Skiena, S.: Deepwalk: Online learning of social representations. In: *Proceedings of the 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 701–710 (2014)
 - [10] Kazemi, B., Abhari, A.: Content-based node2vec for representation of papers in the scientific literature. *Data & Knowledge Engineering* **127**, 101794 (2020)
 - [11] Feldman, R., Golumbic, M.C.: Optimization algorithms for student scheduling via constraint satisfaction. *The Computer Journal* **33**(4), 356–364 (1990)
 - [12] Qi, X.: Weighted jump in random walk graph sampling. *Neurocomputing* **586**, 127581 (2024)
 - [13] Cen, K.-T., Shen, H.-W., Cao, Q., Xu, B.-B., Cheng, X.-Q.: Identity-preserving adversarial training for robust network embedding. *Journal of Computer Science and Technology* **39**(1), 177–191 (2024) <https://doi.org/10.1007/s11390-023-2256-4>
 - [14] Niu, S., Zhou, D.: Soop: Efficient distributed graph computation supporting second-order random walks. *Journal of Computer Science and Technology* **36**(5), 985–1001 (2021) <https://doi.org/10.1007/s11390-021-1234-y>
 - [15] Yang, J., Sheng, Y.-Q., Wang, J.-L., Ni, H.: Cagcn: Centrality-aware graph convolution network for anomaly detection in industrial control systems. *Journal of Computer Science and Technology* **39**(4), 967–983 (2024) <https://doi.org/10.1007/s11390-022-2149-y>
 - [16] Mikolov, T., Chen, K., Corrado, G., Dean, J.: Efficient estimation of word representations in vector space. In: *1st International Conference on Learning Representations, ICLR 2013, Scottsdale, Arizona, USA, May 2-4, 2013, Workshop Track Proceedings* (2013)
 - [17] Zaki, M.J., Meira, W., Meira, W.: *Data Mining and Machine Learning: Fundamental Concepts and Algorithms*. Cambridge University Press, ??? (2020)